

# ercedes-benze-sangini-phansekar

January 30, 2024

```
[1]: #Importing Libraries
```

```
[22]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.feature_selection import VarianceThreshold
variance = VarianceThreshold(threshold=0)
from sklearn.preprocessing import StandardScaler, LabelEncoder
label = LabelEncoder
```

```
[23]: #Reading Dataset
```

```
[24]: train = pd.read_csv('train.csv')
train.head()
```

```
[24]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 378 columns]

```
[25]: test = pd.read_csv('test.csv')
test.head()
```

```
[25]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	\
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

```
[26]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

```
[27]: test.describe()
```

```
[27]:
```

	ID	X10	X11	X12	X13	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	4211.039202	0.019007	0.000238	0.074364	0.061060	
std	2423.078926	0.136565	0.015414	0.262394	0.239468	
min	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	2115.000000	0.000000	0.000000	0.000000	0.000000	
50%	4202.000000	0.000000	0.000000	0.000000	0.000000	
75%	6310.000000	0.000000	0.000000	0.000000	0.000000	
max	8416.000000	1.000000	1.000000	1.000000	1.000000	

	X14	X15	X16	X17	X18	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.427893	0.000713	0.002613	0.008791	0.010216	...	
std	0.494832	0.026691	0.051061	0.093357	0.100570	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	1.000000	0.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X375	X376	X377	X378	X379	\
--	------	------	------	------	------	---

count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.325968	0.049656	0.311951	0.019244	0.011879
std	0.468791	0.217258	0.463345	0.137399	0.108356
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.008791	0.000475	0.000713	0.001663
std	0.089524	0.093357	0.021796	0.026691	0.040752
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 369 columns]

```
[28]: train.isnull().sum()
```

```
[28]: ID      0
      y      0
      X0      0
      X1      0
      X2      0
      ..
      X380    0
      X382    0
      X383    0
      X384    0
      X385    0
      Length: 378, dtype: int64
```

```
[29]: train_target = train["y"]
      train_data = train.drop(["y","ID"],axis=1)
```

```
[30]: train_data.head()
```

```
[30]:   X0 X1  X2 X3 X4 X5 X6 X8  X10 X11  ...  X375  X376  X377  X378  X379  \
0   k  v  at  a  d  u  j  o    0    0  ...    0    0    1    0    0
1   k  t  av  e  d  y  l  o    0    0  ...    1    0    0    0    0
2  az  w   n  c  d  x  j  x    0    0  ...    0    0    0    0    0
3  az  t   n  f  d  x  l  e    0    0  ...    0    0    0    0    0
4  az  v   n  f  d  h  d  n    0    0  ...    0    0    0    0    0
```

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 376 columns]

```
[31]: #To check Variance
      #If for any columns, the variance is equal to zero, then you need to remove
      ↪ those variables
```

```
[32]: train_data.var().sort_values().head(15)
```

/tmp/ipykernel\_162/2491115096.py:1: FutureWarning: The default value of numeric\_only in DataFrame.var is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
train_data.var().sort_values().head(15)
```

```
[32]: X233      0.000000
      X11      0.000000
      X107     0.000000
      X293     0.000000
      X290     0.000000
      X330     0.000000
      X93      0.000000
      X268     0.000000
      X297     0.000000
      X235     0.000000
      X347     0.000000
      X289     0.000000
      X207     0.000238
      X257     0.000238
      X39      0.000238
      dtype: float64
```

```
[33]: test = pd.read_csv('test.csv')
      test.head()
```

```
[33]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0    0    0    1    0    0
1    2   t  b  ai  a  d  b  g  y    0  ...    0    0    1    0    0    0
2    3  az  v  as  f  d  a  j  j    0  ...    0    0    0    1    0    0
```

3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

[34]: test.describe()

[34]:

	ID	X10	X11	X12	X13	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	4211.039202	0.019007	0.000238	0.074364	0.061060	
std	2423.078926	0.136565	0.015414	0.262394	0.239468	
min	1.000000	0.000000	0.000000	0.000000	0.000000	
25%	2115.000000	0.000000	0.000000	0.000000	0.000000	
50%	4202.000000	0.000000	0.000000	0.000000	0.000000	
75%	6310.000000	0.000000	0.000000	0.000000	0.000000	
max	8416.000000	1.000000	1.000000	1.000000	1.000000	

	X14	X15	X16	X17	X18	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.427893	0.000713	0.002613	0.008791	0.010216	...	
std	0.494832	0.026691	0.051061	0.093357	0.100570	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	1.000000	0.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

	X375	X376	X377	X378	X379	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	0.325968	0.049656	0.311951	0.019244	0.011879	
std	0.468791	0.217258	0.463345	0.137399	0.108356	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	1.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.008791	0.000475	0.000713	0.001663

std	0.089524	0.093357	0.021796	0.026691	0.040752
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 369 columns]

```
[35]: train.isnull().sum()
```

```
[35]: ID      0
      y      0
      X0      0
      X1      0
      X2      0
      ..
      X380    0
      X382    0
      X383    0
      X384    0
      X385    0
      Length: 378, dtype: int64
```

```
[36]: train_data_without_zero_var = variance.fit_transform(train_data.iloc[:,9:])
      train_data_without_zero_var
```

```
[36]: array([[0, 1, 0, ..., 0, 0, 0],
              [0, 0, 0, ..., 0, 0, 0],
              [0, 0, 0, ..., 0, 0, 0],
              ...,
              [1, 1, 0, ..., 0, 0, 0],
              [0, 0, 1, ..., 0, 0, 0],
              [0, 0, 0, ..., 0, 0, 0]])
```

```
[37]: labeled_data = train_data.iloc[:,0:8]
      labeled_data.head()
```

```
[37]:   X0 X1  X2 X3 X4 X5 X6 X8
0    k v  at  a  d  u  j  o
1    k t  av  e  d  y  l  o
2  az  w   n  c  d  x  j  x
3  az  t   n  f  d  x  l  e
4  az  v   n  f  d  h  d  n
```

```
[38]: # Checking Null and Unique values for train
```

```
[39]: labeled_data.nunique()
```

```
[39]: X0      47
      X1      27
      X2      44
      X3       7
      X4       4
      X5      29
      X6      12
      X8      25
      dtype: int64
```

```
[40]: labeled_data1 = labeled_data.apply(label().fit_transform)
      labeled_data1.head()
```

```
[40]:   X0  X1  X2  X3  X4  X5  X6  X8
0   32  23  17   0   3  24   9  14
1   32  21  19   4   3  28  11  14
2   20  24  34   2   3  27   9  23
3   20  21  34   5   3  27  11   4
4   20  23  34   5   3  12   3  13
```

```
[41]: labeled_data1.var()
```

```
[41]: X0      188.741938
      X1       72.777974
      X2     118.808135
      X3       3.027295
      X4       0.005461
      X5      68.076236
      X6       8.508730
      X8      49.531868
      dtype: float64
```

```
[42]: train_data_Zero_var_final = pd.DataFrame(train_data_without_zero_var)
      train_data_Zero_var_final.head()
```

```
[42]:   0    1    2    3    4    5    6    7    8    9    ...  345  346  347  348  \
0   0    1    0    0    0    0    1    0    0    1    ...   0    0    1    0
1   0    0    0    0    0    0    1    0    0    0    ...   1    0    0    0
2   0    0    0    0    0    1    0    0    0    0    ...   0    0    0    0
3   0    0    0    0    0    0    0    0    0    0    ...   0    0    0    0
4   0    0    0    0    0    0    0    0    0    0    ...   0    0    0    0

      349  350  351  352  353  354
0   0    0    0    0    0    0
1   0    0    0    0    0    0
```

```

2    0    0    1    0    0    0
3    0    0    0    0    0    0
4    0    0    0    0    0    0

```

[5 rows x 355 columns]

```
[43]: final_train_data= pd.concat([labeled_data1,train_data_Zero_var_final],axis=1)
final_train_data.head()
```

```
[43]:   X0  X1  X2  X3  X4  X5  X6  X8  0  1  ...  345  346  347  348  349  350  \
0   32  23  17   0   3  24   9  14  0  1  ...    0    0    1    0    0    0
1   32  21  19   4   3  28  11  14  0  0  ...    1    0    0    0    0    0
2   20  24  34   2   3  27   9  23  0  0  ...    0    0    0    0    0    0
3   20  21  34   5   3  27  11   4  0  0  ...    0    0    0    0    0    0
4   20  23  34   5   3  12   3  13  0  0  ...    0    0    0    0    0    0

```

```

      351  352  353  354
0      0    0    0    0
1      0    0    0    0
2      1    0    0    0
3      0    0    0    0
4      0    0    0    0

```

[5 rows x 363 columns]

```
[44]: final_train_data.isnull().any()
```

```
[44]: X0      False
X1      False
X2      False
X3      False
X4      False
...
350     False
351     False
352     False
353     False
354     False
Length: 363, dtype: bool

```

```
[45]: test = test.drop(['ID'],axis=1)
test.head()
```

```
[45]:   X0 X1  X2 X3 X4 X5 X6 X8  X10 X11  ...  X375  X376  X377  X378  X379  \
0  az  v   n  f  d  t  a  w    0    0  ...    0    0    0    1    0
1   t  b  ai  a  d  b  g  y    0    0  ...    0    0    1    0    0
2  az  v  as  f  d  a  j  j    0    0  ...    0    0    0    1    0

```



3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0

	X380	X382	X383	X384	X385
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

[5 rows x 376 columns]

```
[46]: #check for null and unique values for test
```

```
[47]: test.nunique()
```

```
[47]: X0      49
      X1      27
      X2      45
      X3       7
      X4       4
      ..
      X380     2
      X382     2
      X383     2
      X384     2
      X385     2
      Length: 376, dtype: int64
```

```
[48]: test.isnull().any()
```

```
[48]: X0      False
      X1      False
      X2      False
      X3      False
      X4      False
      ...
      X380    False
      X382    False
      X383    False
      X384    False
      X385    False
      Length: 376, dtype: bool
```

```
[49]: test.var().sort_values().head(15)
```

/tmp/ipykernel\_162/1038450595.py:1: FutureWarning: The default value of

numeric\_only in DataFrame.var is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
test.var().sort_values().head(15)
```

```
[49]: X369    0.000000
      X257    0.000000
      X258    0.000000
      X296    0.000000
      X295    0.000000
      X288    0.000238
      X330    0.000238
      X210    0.000238
      X290    0.000238
      X293    0.000238
      X259    0.000238
      X11     0.000238
      X105    0.000238
      X278    0.000238
      X233    0.000238
      dtype: float64
```

```
[50]: test_without_zero_var=variance.transform(test.iloc[:,9:])
      test_without_zero_var
```

```
[50]: array([[0, 0, 0, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0],
            [0, 0, 1, ..., 0, 0, 0],
            ...,
            [0, 0, 1, ..., 0, 0, 0],
            [0, 1, 1, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0]])
```

```
[51]: test_without_zero_var_final = pd.DataFrame(test_without_zero_var)
      test_without_zero_var_final
```

```
[51]:
```

	0	1	2	3	4	5	6	7	8	9	...	345	346	347	\
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
1	0	0	0	0	0	0	0	1	0	0	...	0	0	1	
2	0	0	1	0	0	0	0	0	0	0	...	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
4	0	0	1	0	0	0	0	0	0	0	...	1	0	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
4204	0	0	1	0	0	0	0	0	0	0	...	0	0	0	
4205	0	0	0	0	0	0	0	1	0	0	...	0	1	0	
4206	0	0	1	0	0	0	0	0	0	0	...	0	0	0	

4207	0	1	1	0	0	0	0	0	0	0	...	0	0	1
4208	0	0	0	0	0	0	0	1	0	0	...	1	0	0

	348	349	350	351	352	353	354
0	1	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0
3	1	0	0	0	0	0	0
4	0	0	0	0	0	0	0
...	...	...	...	...	...	...	...
4204	0	0	0	0	0	0	0
4205	0	0	0	0	0	0	0
4206	0	0	0	0	0	0	0
4207	0	0	0	0	0	0	0
4208	0	0	0	0	0	0	0

[4209 rows x 355 columns]

```
[52]: labeled_data1=test.iloc[:,0:8]
labeled_data1.head()
```

```
[52]:   X0 X1  X2 X3 X4 X5 X6 X8
0  az  v   n  f  d  t  a  w
1   t  b  ai  a  d  b  g  y
2  az  v  as  f  d  a  j  j
3  az  l   n  f  d  z  l  n
4   w  s  as  c  d  y  i  m
```

```
[53]: test_label = labeled_data1.apply(label().fit_transform)
test_label.head()
```

```
[53]:   X0 X1 X2 X3 X4 X5 X6 X8
0  21 23 34  5  3 26  0 22
1  42  3  8  0  3  9  6 24
2  21 23 17  5  3  0  9  9
3  21 13 34  5  3 31 11 13
4  45 20 17  2  3 30  8 12
```

```
[54]: test_data_final = pd.concat([test_label,test_without_zero_var_final],axis=1)
test_data_final.head()
```

```
[54]:   X0 X1 X2 X3 X4 X5 X6 X8 0 1 ... 345 346 347 348 349 350 \
0  21 23 34  5  3 26  0 22 0 0 ...  0  0  0  1  0  0
1  42  3  8  0  3  9  6 24 0 0 ...  0  0  1  0  0  0
2  21 23 17  5  3  0  9  9 0 0 ...  0  0  0  1  0  0
3  21 13 34  5  3 31 11 13 0 0 ...  0  0  0  1  0  0
4  45 20 17  2  3 30  8 12 0 0 ...  1  0  0  0  0  0
```

	351	352	353	354
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 363 columns]

```
[55]: #Perform Dimensionality Reduction
```

```
[56]: test_data_final.keys()
```

```
[56]: Index(['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8',    0,    1,
          ...,
          345, 346, 347, 348, 349, 350, 351, 352, 353, 354],
          dtype='object', length=363)
```

```
[57]: test_data_final.columns=test_data_final.columns.astype(str)
```

```
[58]: scaler = StandardScaler()
      scaler.fit(test_data_final)
```

```
[58]: StandardScaler()
```

```
[59]: scaled_data = scaler.transform(test_data_final)
```

```
[60]: scaled_data
```

```
[60]: array([[ -0.62521149,  1.39576032,  1.58606761, ..., -0.02180363,
          -0.02670705, -0.04081511],
          [ 0.75460919, -0.94519929, -0.95644521, ..., -0.02180363,
          -0.02670705, -0.04081511],
          [-0.62521149,  1.39576032, -0.07634462, ..., -0.02180363,
          -0.02670705, -0.04081511],
          ...,
          [ 1.08313793,  1.39576032, -0.07634462, ..., -0.02180363,
          -0.02670705, -0.04081511],
          [-1.54509194,  1.39576032, -0.07634462, ..., -0.02180363,
          -0.02670705, -0.04081511],
          [ 0.75460919, -1.17929525, -0.95644521, ..., -0.02180363,
          -0.02670705, -0.04081511]])
```

```
[61]: #from sklearn.decomposition import PCA
```

```
[62]: from sklearn.model_selection import train_test_split
```

```
[63]: x_train,x_test, y_train, y_test = train_test_split(scaled_data, train_target,
↳ random_state = 42, test_size = 0.3)
```

```
[64]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
[64]: ((2946, 363), (1263, 363), (2946,), (1263,))
```

```
[65]: from sklearn.decomposition import PCA
pca = PCA(n_components=2)
```

```
[66]: pca=PCA(n_components=2)
pca.fit(scaled_data)
x_pca = pca.transform(scaled_data)
scaled_data.shape
x_pca.shape
```

```
[66]: (4209, 2)
```

```
[67]: scaled_data.shape
```

```
[67]: (4209, 363)
```


```
[68]: x_pca.shape
```

```
[68]: (4209, 2)
```

```
[69]: # plot figure
```

```
[70]: plt.figure(figsize=(8,6))
plt.scatter(x_pca[:,0],x_pca[:,1],c=train['y'],cmap='plasma')
plt.xlabel('First principal Component')
plt.ylabel('Second Principal Component')
```

```
[70]: Text(0, 0.5, 'Second Principal Component')
```



```
[71]: #XGBoost
```

```
[72]: from sklearn import svm
from sklearn.metrics import r2_score, mean_squared_error
from xgboost import XGBRegressor
xgbr = XGBRegressor(random_state=42)
```

```
[73]: model = xgbr.fit(x_train, y_train)
```

```
[74]: ypred_test = model.predict(x_test)
ypred_test
```

```
[74]: array([107.663704, 102.37494 , 102.57482 , ...,  89.09564 , 102.00936 ,
          99.39917 ], dtype=float32)
```

```
[75]: ypred_train = model.predict(x_train)
ypred_train
```

```
[75]: array([ 99.53446 , 102.702675,  96.319984, ...,  98.61626 ,  96.17458 ,
          101.20219 ], dtype=float32)
```

```
[76]: print(r2_score(ypred_train, y_train))
```

0.24660149892560312

```
[77]: print(mean_squared_error(ypred_train, y_train))
```


47.12593276125439

```
[78]: test_data_final_prediction = model.predict(test_data_final)
test_data_final_prediction
```

```
[78]: array([ 82.54247 ,  79.10698 , 105.050095, ..., 102.59405 ,  98.634445,
          109.8387  ], dtype=float32)
```

```
[79]: prediction = pd.DataFrame({'ytest':y_test,'ypred':ypred_test})
```

```
[80]: #sns.lmplot(x=ytest, y=ypred, data=prediction)
plt.plot(prediction['ytest'],color='red')
plt.plot(prediction['ypred'],color='Blue')
plt.show()
```



```
[ ]:
```

```
[ ]:
```